

# Diffusion Model

Computer Vision - Project

---

Clément, Grégoire, Nathan

January 14, 2025

# What is Diffusion ?



Diffusion is the state-of-the-art generative process to generate images for creating new images near the original one. It also allows to generate images from text.

## **First generation: Denoising Diffusion Probabilistic Models**

---

---

## Denoising Diffusion Probabilistic Models

---


**Jonathan Ho**  
UC Berkeley  
`jonathanho@berkeley.edu`

**Ajay Jain**  
UC Berkeley  
`ajayj@berkeley.edu`

**Pieter Abbeel**  
UC Berkeley  
`pabbeel@cs.berkeley.edu`

# General Idea

Consider the set of **hand-written digits**  $D$ . Can you give a probability distribution  $q$  such that  $x \sim q(x)$  ?




**Figure 1:** Source: ludwig.ai

# General Idea


---

Consider the set of **hand-written digits**  $D$ . It is hard to find  $q$  such that  $x \sim q(x)$ , we need a **clever way to sample hand-written digits**. Consider the following process:




# General Idea

Consider the set of **hand-written digits**  $D$ . It is hard to find  $q$  such that  $x \sim q(x)$ , we need a **clever way to sample hand-written digits**. Consider the following process:



# General Idea


Consider the set of **hand-written digits**  $D$ . It is hard to find  $q$  such that  $x \sim q(x)$ , we need a **clever way to sample hand-written digits**. Consider the following process:





# General Idea


Consider the set of **hand-written digits**  $D$ . It is hard to find  $q$  such that  $x \sim q(x)$ , we need a **clever way to sample hand-written digits**. Consider the following process:



# General Idea

Consider the set of **hand-written digits**  $D$ . It is hard to find  $q$  such that  $x \sim q(x)$ , we need a **clever way to sample hand-written digits**.


Consider the following process:



# General Idea

Consider the set of **hand-written digits**  $D$ . It is hard to find  $q$  such that  $x \sim q(x)$ , we need a **clever way to sample hand-written digits**.


Consider the following process:



# General Idea

Consider the set of **hand-written digits**  $D$ . It is hard to find  $q$  such that  $x \sim q(x)$ , we need a **clever way to sample hand-written digits**.


Consider the following process:



Formally:  $q(x_{t+1} | x_t) := \mathcal{N}(x_{t+1}; \sqrt{1 - \beta_t}x_t, \beta_t I)$  for some schedule  $(\beta_t)_t$ . Can we **learn to reverse this process** ?


# What we want to learn

Given a noisy image  $x_t$ , we train a model to predict  $x_{t-1}$ .



# What we want to learn

Given a noisy image  $x_t$ , we train a model to predict  $x_{t-1}$ .



- Given a noisy image  $x_t$  and  $t$ , we sample according to  $p_\theta(x_{t-1} | x_t) := \mathcal{N}(x_{t-1}; \mu_\theta(x_t, t), \Sigma_\theta(x_t, t))$ .

# Decreasing training data generation cost

Given a data image  $x_0$ , computing  $x_t$  takes  $t$  samplings on  $q$ . But a simple trick allows us to do it in only one step.

Remember that  $q(x_{t+1} | x_t) := \mathcal{N}(x_{t+1}; \sqrt{1 - \beta_t}x_t, \beta_t I)$ . Let  $\alpha_t = 1 - \beta_t$  and  $\bar{\alpha}_t = \prod_{i=1}^t \alpha_i$ .

$$\begin{aligned}x_t &= \sqrt{\alpha_t}x_{t-1} + \sqrt{1 - \alpha_t}\epsilon_{t-1} \\&= \sqrt{\alpha_t}\sqrt{\alpha_{t-1}}x_{t-2} + \sqrt{\alpha_t}\sqrt{1 - \alpha_{t-1}}\epsilon_{t-1} + \sqrt{1 - \alpha_t}\epsilon_{t-1} \\&= \sqrt{\alpha_t\alpha_{t-1}}x_{t-2} + \sqrt{\alpha_t(1 - \alpha_{t-1}) + (1 - \alpha_t)}\bar{\epsilon}_t \\&= \sqrt{\alpha_t\alpha_{t-1}}x_{t-2} + \sqrt{1 - \alpha_t\alpha_{t-1}}\bar{\epsilon}_t\end{aligned}$$

Let  $G_1 \sim \mathcal{N}(0, \sigma_1^2 I)$ ,  $G_2 \sim \mathcal{N}(0, \sigma_2^2 I)$ , the sum of them gives  $g_2 \sim \mathcal{N}(0, (\sigma_1^2 + \sigma_2^2)I)$ .

# Decreasing training data generation cost

---

Given a data image  $x_0$ , computing  $x_t$  takes  $t$  samplings on  $q$ . But a simple trick allows us to do it in only one step.

We have  $x_t = \sqrt{\bar{\alpha}_t}x_0 + \sqrt{1 - \bar{\alpha}_t}\epsilon$ .



# Training

---

For now, our model is learning  $\mu$  and  $\Sigma$ , i.e. we sample according to

$$p_{\theta}(x_{t-1} \mid x_t) := \mathcal{N}(x_{t-1}; \mu_{\theta}(x_t, t), \Sigma_{\theta}(x_t, t))$$

They fixed **fixing**  $\Sigma_{\theta}$  to a constant to simplify computations. So,

$$p_{\theta}(x_{t-1} \mid x_t) := \mathcal{N}(x_{t-1}; \mu_{\theta}(x_t, t), \sigma_t I)$$

# Training

---

$$p_{\theta}(x_{t-1} \mid x_t) := \mathcal{N}(x_{t-1}; \mu_{\theta}(x_t, t), \sigma_t I)$$

Using negative log likelihood, approximations and computations, we want to minimize:

$$E_q \left[ \frac{1}{2\sigma_t^2} \|\tilde{\mu}_t(x_t, x_0) - \mu_{\theta}(x_t, t)\|^2 \right]$$

where  $\tilde{\mu}$  is the optimal mean that depends on  $x_0$ .

# Training

$$p_{\theta}(x_{t-1} \mid x_t) := \mathcal{N}(x_{t-1}; \mu_{\theta}(x_t, t), \sigma_t I)$$

Using negative log likelihood, approximations and computations, we want to minimize:


$$E_q \left[ \frac{1}{2\sigma_t^2} \|\tilde{\mu}_t(x_t, x_0) - \mu_{\theta}(x_t, t)\|^2 \right]$$

where  $\tilde{\mu}$  is the optimal mean that depends on  $x_0$  which we don't know. We can also have  $\epsilon$  using

$$E_q \left[ \frac{\beta_t^2}{2\sigma_t^2 \alpha_t (1 - \bar{\alpha}_t)} \|\epsilon - \epsilon_{\theta}(\sqrt{\bar{\alpha}_t} x_0 + \sqrt{1 - \bar{\alpha}_t} \epsilon, t)\|^2 \right]$$


# Our results - Gaussian

We have started with Gaussian generation:




# Our results - Gaussian

We have started with Gaussian generation and got satisfying results:




# Our results - Spirale

Then we moved to a more complicated dataset, Spirale generation:




# Our results - Spirale

Then we moved to a more complicated dataset, **Spirale generation** and also got satisfying results:



# Generating Images


We need a model powerful enough to predict  $\epsilon$  (or  $\mu$ ) for 2D images (the size of  $\epsilon$ ), for this we use UNet.





# Our results - MNIST Generation

Generated Images



# Amelioration

---

## Auto-Encoding Variational Bayes

---

**Diederik P. Kingma**  
Machine Learning Group  
Universiteit van Amsterdam  
dpkingma@gmail.com

**Max Welling**  
Machine Learning Group  
Universiteit van Amsterdam  
welling.max@gmail.com

Imagine that you have  $T = 10,000$ . You can not do 10,000 steps to generate one image. But, using the reparametrization trick:

$$x_T = x_0 + \sigma_{T,0}\epsilon_T \Rightarrow \tilde{x}_0$$
$$x_{T-k} = \tilde{x}_0 + \sigma_{T-k,0}\epsilon_{T-k}$$

So we can generate with  $T/k$  steps.

# OpenAI's incrementation

---

---

## Improved Denoising Diffusion Probabilistic Models

---

Alex Nichol<sup>\*1</sup> Prafulla Dhariwal<sup>\*1</sup>

# OpenAI's incrementation

---

## Improved Denoising Diffusion Probabilistic Models

---

Alex Nichol<sup>\*1</sup> Prafulla Dhariwal<sup>\*1</sup>

Recall that:

- $\Sigma_{\theta}$  was set to a constant,

# OpenAI's incrementation

---

## Improved Denoising Diffusion Probabilistic Models

---

Alex Nichol<sup>\*1</sup> Prafulla Dhariwal<sup>\*1</sup>

Recall that:

- $\Sigma_{\theta}$  was set to a constant,
- The training loss was a simplified version of reality.

# OpenAI's incrementation

---

## Improved Denoising Diffusion Probabilistic Models

---

Alex Nichol<sup>\*1</sup> Prafulla Dhariwal<sup>\*1</sup>

Recall that:

- $\Sigma_{\theta}$  was set to a constant,
- The training loss was a simplified version of reality.

This paper tackle these problems.

# How to modelize $\Sigma_\theta$

---

Ho et al. fixed  $\Sigma_\theta(x_t, t) = \sigma_t^2 I$  with:

- $\sigma_t^2 = \beta_t$  optimal if  $x_0 \sim \mathcal{N}(0, I)$ ,



# How to modelize $\Sigma_\theta$

---

Ho et al. fixed  $\Sigma_\theta(x_t, t) = \sigma_t^2 I$  with:

- $\sigma_t^2 = \beta_t$  optimal if  $x_0 \sim \mathcal{N}(0, I)$ ,
- $\sigma_t^2 = \bar{\beta}_t$  optimal if  $x_0$  is a point.

# How to modelize $\Sigma_\theta$

Ho et al. fixed  $\Sigma_\theta(x_t, t) = \sigma_t^2 I$  with:

- $\sigma_t^2 = \beta_t$  optimal if  $x_0 \sim \mathcal{N}(0, I)$ ,
- $\sigma_t^2 = \bar{\beta}_t$  optimal if  $x_0$  is a point.

Hence for each  $t$ , we know that  $\Sigma^*(x, t)$  is between  $\beta_t$  and  $\bar{\beta}_t$ .

# How to modelize $\Sigma_\theta$

Ho et al. fixed  $\Sigma_\theta(x_t, t) = \sigma_t^2 I$  with:

- $\sigma_t^2 = \beta_t$  optimal if  $x_0 \sim \mathcal{N}(0, I)$ ,
- $\sigma_t^2 = \bar{\beta}_t$  optimal if  $x_0$  is a point.

Hence for each  $t$ , we know that  $\Sigma^*(x, t)$  is between  $\beta_t$  and  $\bar{\beta}_t$ .

Ho et al. have found that the impact is negligible...

# How to modelize $\Sigma_\theta$

Hence for each  $t$ , we know that  $\Sigma^*(x, t)$  is between  $\beta_t$  and  $\bar{\beta}_t$ .

Ho et al. have found that the impact is negligible...




Figure 1. The ratio  $\tilde{\beta}_t/\beta_t$  for every diffusion step for diffusion processes of different lengths.

# How to modelize $\Sigma_\theta$

Hence for each  $t$ , we know that  $\Sigma^*(x, t)$  is between  $\beta_t$  and  $\bar{\beta}_t$ .

Ho et al. have found that the impact is negligible. But it depends of other hyperparameters.

Hence, they interpolate between the two extreme values, and let the model learn  $v(t)$ :

$$\Sigma_\theta(x_t, t) := \exp(v_\theta(t) \log(\beta_t) + (1 - v_\theta(t)) \log(\bar{\beta}_t))$$

# Changing the Loss

---

In the paper of Ho, they use the loss  $L_{simple}$ , a simplified version of  $L_{LVB}$ , that only affects  $\mu_\theta$ .

# Changing the Loss

---

In the paper of Ho, they use the loss  $L_{simple}$ , a simplified version of  $L_{LVB}$ , that only affects  $\mu_{\theta}$ .

Optimizing  $L_{simple}$  performs well for FID, but poorly for log-likelihood metric.

# Changing the Loss

---

In the paper of Ho, they use the loss  $L_{simple}$ , a simplified version of  $L_{LVB}$ , that only affects  $\mu_\theta$ .

Optimizing  $L_{simple}$  performs well for FID, but poorly for log-likelihood metric.

$L_{LVB}$  affects  $\Sigma_\theta$  but increasing log-likelihood score can increase FID, hence we need a trade-of.



# Changing the Loss

In the paper of Ho, they use the loss  $L_{simple}$ , a simplified version of  $L_{LVB}$ , that only affects  $\mu_\theta$ .

Optimizing  $L_{simple}$  performs well for FID, but poorly for log-likelihood metric.

$L_{LVB}$  affects  $\Sigma_\theta$  but increasing log-likelihood score can increase FID, hence we need a trade-of.

They train on

$$L = L_{simple} + \lambda L_{LVB}$$

# Changing the Loss

In the paper of Ho, they use the loss  $L_{simple}$ , a simplified version of  $L_{LVB}$ , that only affects  $\mu_\theta$ .

Optimizing  $L_{simple}$  performs well for FID, but poorly for log-likelihood metric.

$L_{LVB}$  affects  $\Sigma_\theta$  but increasing log-likelihood score can increase FID, hence we need a trade-of.

They train on

$$L = L_{simple} + \lambda L_{LVB}$$


This loss is prone to gradient exploding and we need importance sampling to implement it.

## **Classifier Guidance**

---

# Importance of labels


Let's get back to hand-written digits generation:



**Figure 1:** Source: ludwig.ai

# Importance of labels

Let's get back to hand-written digits generation:



**Figure 1:** Source: ludwig.ai

A DDPM can generate new images that look like digits, but the model can't distinguish a mix of two digits and a real digit.

# Importance of labels

---

Let's get back to hand-written digits generation:

If we have a classifier that gives  $p_\phi(y \mid x_t)$ , we can sample using

$$p_{\theta,\phi}(x_t \mid x_{t+1}, y) = Z p_\theta(x_t \mid x_{t+1}) p_\phi(y \mid x_t)$$

This way if we set  $y = 3$ , we can trick our model to generate something that looks like a 3.


# Langevin Dynamics

---

Given  $x_0 \sim \pi(x)$  an unknown distribution, if we iterate through  $x_{i+1} \leftarrow x_i + \epsilon \nabla_x \log p(x) + \sqrt{2\epsilon} z_i$  with  $\epsilon \rightarrow 0$  and  $z_i \sim \mathcal{N}(0, I)$ , we can sample from  $p(x)$ .

# Langevin Dynamics

Given  $x_0 \sim \pi(x)$  an unknown distribution, if we iterate through  $x_{i+1} \leftarrow x_i + \epsilon \nabla_x \log p(x) + \sqrt{2\epsilon} z_i$  with  $\epsilon \rightarrow 0$  and  $z_i \sim \mathcal{N}(0, I)$ , we can sample from  $p(x)$ .




**Figure 1:** Visualizations from Yang Song's work.



# Langevin Dynamics


Given  $x_0 \sim \pi(x)$  an unknown distribution, if we iterate through  $x_{i+1} \leftarrow x_i + \epsilon \nabla_x \log p(x) + \sqrt{2\epsilon} z_i$  with  $\epsilon \rightarrow 0$  and  $z_i \sim \mathcal{N}(0, I)$ , we can sample from  $p(x)$ .



**Figure 1:** Visualizations from Yang Song's work.

# Langevin Dynamics


Given  $x_0 \sim \pi(x)$  an unknown distribution, if we iterate through  $x_{i+1} \leftarrow x_i + \epsilon \nabla_x \log p(x) + \sqrt{2\epsilon} z_i$  with  $\epsilon \rightarrow 0$  and  $z_i \sim \mathcal{N}(0, I)$ , we can sample from  $p(x)$ .



**Figure 1:** Visualizations from Yang Song's work.

# Langevin Dynamics

Given  $x_0 \sim \pi(x)$  an unknown distribution, if we iterate through  $x_{i+1} \leftarrow x_i + \epsilon \nabla_x \log p(x) + \sqrt{2\epsilon} z_i$  with  $\epsilon \rightarrow 0$  and  $z_i \sim \mathcal{N}(0, I)$ , we can sample from  $p(x)$ .



**Figure 1:** Visualizations from Yang Song's work.

# Langevin Dynamics

Given  $x_0 \sim \pi(x)$  an unknown distribution, if we iterate through  $x_{i+1} \leftarrow x_i + \epsilon \nabla_x \log p(x) + \sqrt{2\epsilon} z_i$  with  $\epsilon \rightarrow 0$  and  $z_i \sim \mathcal{N}(0, I)$ , we can sample from  $p(x)$ .

So we need to know  $\nabla_x \log p(x)$ , but don't need to know  $p(x)$ .

# Langevin Dynamics

Given  $x_0 \sim \pi(x)$  an unknown distribution, if we iterate through  $x_{i+1} \leftarrow x_i + \epsilon \nabla_x \log p(x) + \sqrt{2\epsilon} z_i$  with  $\epsilon \rightarrow 0$  and  $z_i \sim \mathcal{N}(0, I)$ , we can sample from  $p(x)$ .

So we need to know  $\nabla_x \log p(x)$ , but don't need to know  $p(x)$ .

From the classifier  $p_\phi$ , one can get an approximation of  $\nabla_x \log p(x)$ .

# Low vs High temperature

---



**Figure 1:** Classifier-free Diffusion Guidance

# Low vs High temperature



**Figure 1:** Classifier-free Diffusion Guidance

- Low-temperature optimizes FID score

# Low vs High temperature



**Figure 1:** Classifier-free Diffusion Guidance

- Low-temperature optimizes FID score
- High-temperature optimizes Inception score



# Low vs High temperature



**Figure 1:** Classifier-free Diffusion Guidance

- Low-temperature optimizes FID score
- High-temperature optimizes Inception score

We can do a trade-off by following more  $p_\phi$  or  $p_\theta$  between exploration and distance to the original distribution.

**GLIDE: draw what you prompt**

---

# Classifier-Free Guidance

---

Previous guidance need a trained classifier. They define:

- An unconditional DDPM, that predicts  $p_{\theta}(z)$ .
- A conditional DPPM, that predicts  $p_{\theta}(z | c)$ .

# Classifier-Free Guidance

Previous guidance need a trained classifier. They define:

- An unconditional DDPM, that predicts  $p_{\theta}(z)$ .
- A conditional DPPM, that predicts  $p_{\theta}(z | c)$ .

Rather than using  $p_{\phi}$ , they train both models simultaneously and use the gradient of  $p_{\theta}(z | c)$  to estimate  $\nabla \log p(z | c)$ .

To force a label  $c$ , we train a model that estimates the probability for  $x$  to be of class  $c$ , and uses its gradient.

To force a label  $c$ , we train a model that estimates the probability for  $x$  to be of class  $c$ , and uses its gradient.

To force a sentence  $c$ , we train a model that estimates the distance between  $x$  and the sentence  $c$ , and uses its gradient.

# CLIP

To force a sentence  $c$ , we train a model that estimates the distance between  $x$  and the sentence  $c$ , and uses its gradient.

## 1. Contrastive pre-training

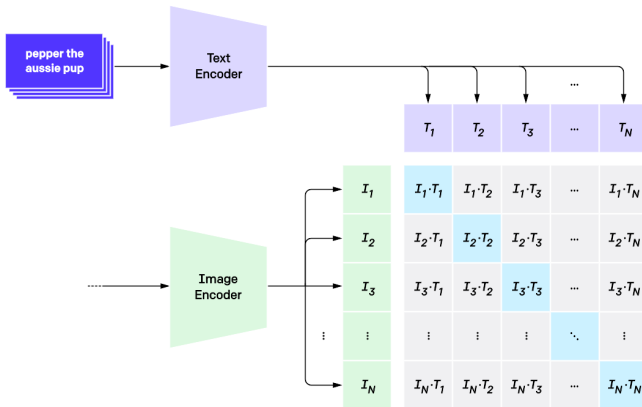


Figure 2: How to compare an image and a text

# Results (OpenAI GLIDE)



"a hedgehog using a calculator"



"a corgi wearing a red bowtie and a purple party hat"



"robots meditating in a vipassana retreat"



"a fall landscape with a small cottage next to a lake"



"a surrealist dream-like oil painting by salvador dali of a cat playing checkers"




"a professional photo of a sunset behind the grand canyon"



"a high-quality oil painting of a psychedelic hamster dragon"



"an illustration of albert einstein wearing a superhero costume"



"a boat in the canals of venice"



"a painting of a fox in the style of starry night"



"a red cube on top of a blue cube"



"a stained glass window of a panda eating bamboo"



## **State-of-the-art brief review**

---

GLIDE **accepts** one more input than text: **a mask for inpainting**.

# ControlNet

---

GLIDE accepts one more input than text: a mask for inpainting.

ControlNet generalises it by adding more optional inputs (e.g. Cany edges representation / Human Pose / Sketch).

# ControlNet

Canny ControlNet model output



Normal Map ControlNet model output

